

Trabalho 3 - Redes Neurais

Fazer um script em jupyter notebook para classificação de imagens.

Você deve plotar as curvas de erro de treino e validação.

Use a metodologia de validação Holdout.

Ao final, você deve gerar métricas para o conjunto de teste.

Métricas de classificação: acurácia e f1 por classe. (consulte o scikit-learn)

Imports

Primeiro é preciso importar as bibliotecas e funções a serem utilizadas.

```
In [1]: import torch
from torchvision import datasets
from torchvision.transforms import ToTensor
import torch.utils.data as data_utils
from torch import nn
import matplotlib.pyplot as plt
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, classifica
```

Pegando os datasets

Utilizando o dataset CIFAR10 que possui link direto com o torchvision. É um dataset constituídos por imagens coloridas de veículos e animais onde existem 10 classes diferentes. O objetivo é criar um classificador a partir de uma CNN.

[Link para o dataset](#)

```
In [2]: train_dataset_base = datasets.CIFAR10(
    root = "data",
    train = True,
    transform = ToTensor(),
)

test_dataset = datasets.CIFAR10(
    root = "data",
    train = False,
    transform = ToTensor()
)

print("Shape of the train_dataset: ", train_dataset_base.data.shape)
print("Shape of the test_dataset: ", test_dataset.data.shape)
print("Classes: ", train_dataset_base.classes)
```

```
Shape of the train_dataset: (50000, 32, 32, 3)
Shape of the test_dataset: (10000, 32, 32, 3)
Classes: ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```

Dividindo e Transformando o Dataset

```
In [3]: # Criando um dataset de validação a partir do treino
train_dataset, valid_dataset = data_utils.random_split(train_dataset_base, [40000, 10000])

# Criando os dataloaders com o batch_size de 64
train_loader = data_utils.DataLoader(train_dataset, batch_size=64, shuffle=True)
valid_loader = data_utils.DataLoader(valid_dataset, batch_size=64, shuffle=True)
test_loader = data_utils.DataLoader(test_dataset, batch_size=64, shuffle=True)
```

Definindo o device

Como minha placa de vídeo é AMD não possuo Cuda.

```
In [4]: device = 'cuda' if torch.cuda.is_available() else 'cpu'
print('Using {} device'.format(device))
```

Using cpu device

Definindo a Rede Neural

Criei uma classe que herda nn.Module e defini as camadas da rede aqui. Usei um modelo base de CNN apenas para facilitar.

```
In [5]: class NeuralNetwork(nn.Module):

    def __init__(self):
        super(NeuralNetwork, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(nn.functional.relu(self.conv1(x)))
        x = self.pool(nn.functional.relu(self.conv2(x)))
        x = x.view(-1, 16 * 5 * 5)
        x = nn.functional.relu(self.fc1(x))
        x = nn.functional.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

Jogando a Rede para a CPU

```
In [6]: model = NeuralNetwork().to(device)
print(model)
```

```

NeuralNetwork(
  (conv1): Conv2d(3, 6, kernel_size=(5, 5), stride=(1, 1))
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (conv2): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))
  (fc1): Linear(in_features=400, out_features=120, bias=True)
  (fc2): Linear(in_features=120, out_features=84, bias=True)
  (fc3): Linear(in_features=84, out_features=10, bias=True)
)

```

Definindo a loss function e o otimizador

Learning rate inicial de 0.01 pois foi a com melhor desempenho.

```

In [7]: loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=0.01)

```

Treinando e Validando a Rede

Funções para implementar o treino e a validação

Realizam os ajustes de pesos das redes e calculam a perda e acurácia do algoritmo

```

In [8]: def train_loop(dataloader, model, loss_fn, optimizer):
    size = len(dataloader.dataset)

    loss_sum = 0
    i = 0
    for X, y in dataloader:
        pred = model(X)
        loss = loss_fn(pred, y)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        loss_sum += loss.item() * len(X)

    avg_loss = loss_sum / size
    print(f"Train Error: Avg loss = {avg_loss:>8f}")
    return avg_loss

```

```

In [9]: def valid_loop(dataloader, model, loss_fn):
    size = len(dataloader.dataset)
    loss, correct = 0, 0

    with torch.no_grad():
        for X, y in dataloader:
            pred = model(X)
            loss += loss_fn(pred, y).item() * len(X)
            correct += (pred.argmax(1) == y).type(torch.float).sum().item()

    avg_loss = loss / size
    acc = correct / size
    print(f"Test Error: Accuracy = {100*acc:>0.1f}%, Avg loss = {avg_loss:>8f}")

```

```
return avg_loss, acc
```

Loop do Treino

Utilizei 50 epochs pois o desempenho não melhorava muito depois disso. Salvei o modelo com melhor desempenho de perda na validação.

```
In [10]: n_epoch = 50
loss_valid_history = []
acc_history = []
loss_train_history = []
min_loss_valid = np.Inf

for epoch in range(n_epoch):
    print(f"Epoch {epoch+1}\n-----")
    loss_train = train_loop(train_loader, model, loss_fn, optimizer)
    loss_valid, accuracy = valid_loop(valid_loader, model, loss_fn)
    if loss_valid <= min_loss_valid:
        print(f"Validation loss decreased ({min_loss_valid:>0.6f} --> {loss_valid:>0.6f})")
        min_loss_valid = loss_valid
        torch.save(model.state_dict(), 'models/best_model_trab3.pt')
    print("-----\n")

    loss_train_history.append(loss_train)
    loss_valid_history.append(loss_valid)
    acc_history.append(accuracy)
```

Epoch 1

```
-----  
Train Error: Avg loss = 2.301983  
Test Error: Accuracy = 12.0%, Avg loss = 2.299056  
Validation loss decreased (inf --> 2.299056). Saving model ...  
-----
```

Epoch 2

```
-----  
Train Error: Avg loss = 2.292651  
Test Error: Accuracy = 16.4%, Avg loss = 2.277575  
Validation loss decreased (2.299056 --> 2.277575). Saving model ...  
-----
```

Epoch 3

```
-----  
Train Error: Avg loss = 2.209371  
Test Error: Accuracy = 23.3%, Avg loss = 2.132837  
Validation loss decreased (2.277575 --> 2.132837). Saving model ...  
-----
```

Epoch 4

```
-----  
Train Error: Avg loss = 2.039506  
Test Error: Accuracy = 28.0%, Avg loss = 1.994714  
Validation loss decreased (2.132837 --> 1.994714). Saving model ...  
-----
```

Epoch 5

```
-----  
Train Error: Avg loss = 1.946053  
Test Error: Accuracy = 30.0%, Avg loss = 1.905836  
Validation loss decreased (1.994714 --> 1.905836). Saving model ...  
-----
```

Epoch 6

```
-----  
Train Error: Avg loss = 1.864871  
Test Error: Accuracy = 34.8%, Avg loss = 1.819797  
Validation loss decreased (1.905836 --> 1.819797). Saving model ...  
-----
```

Epoch 7

```
-----  
Train Error: Avg loss = 1.784232  
Test Error: Accuracy = 37.3%, Avg loss = 1.754910  
Validation loss decreased (1.819797 --> 1.754910). Saving model ...  
-----
```

Epoch 8

```
-----  
Train Error: Avg loss = 1.720086  
Test Error: Accuracy = 39.7%, Avg loss = 1.687446  
Validation loss decreased (1.754910 --> 1.687446). Saving model ...  
-----
```

Epoch 9

```
-----  
Train Error: Avg loss = 1.655481  
Test Error: Accuracy = 42.1%, Avg loss = 1.631101
```

Validation loss decreased (1.687446 --> 1.631101). Saving model ...

Epoch 10

Train Error: Avg loss = 1.604435

Test Error: Accuracy = 43.7%, Avg loss = 1.574025

Validation loss decreased (1.631101 --> 1.574025). Saving model ...

Epoch 11

Train Error: Avg loss = 1.560346

Test Error: Accuracy = 45.2%, Avg loss = 1.545079

Validation loss decreased (1.574025 --> 1.545079). Saving model ...

Epoch 12

Train Error: Avg loss = 1.521284

Test Error: Accuracy = 45.1%, Avg loss = 1.538583

Validation loss decreased (1.545079 --> 1.538583). Saving model ...

Epoch 13

Train Error: Avg loss = 1.493416

Test Error: Accuracy = 47.3%, Avg loss = 1.475665

Validation loss decreased (1.538583 --> 1.475665). Saving model ...

Epoch 14

Train Error: Avg loss = 1.462759

Test Error: Accuracy = 48.1%, Avg loss = 1.459460

Validation loss decreased (1.475665 --> 1.459460). Saving model ...

Epoch 15

Train Error: Avg loss = 1.436622

Test Error: Accuracy = 48.1%, Avg loss = 1.457191

Validation loss decreased (1.459460 --> 1.457191). Saving model ...

Epoch 16

Train Error: Avg loss = 1.412990

Test Error: Accuracy = 48.9%, Avg loss = 1.432100

Validation loss decreased (1.457191 --> 1.432100). Saving model ...

Epoch 17

Train Error: Avg loss = 1.393898

Test Error: Accuracy = 48.6%, Avg loss = 1.437844

Epoch 18

```
Train Error: Avg loss = 1.371760
Test Error: Accuracy = 50.9%, Avg loss = 1.391630
Validation loss decreased (1.432100 --> 1.391630). Saving model ...
-----

Epoch 19
-----
Train Error: Avg loss = 1.354982
Test Error: Accuracy = 50.1%, Avg loss = 1.393070
-----

Epoch 20
-----
Train Error: Avg loss = 1.336949
Test Error: Accuracy = 51.7%, Avg loss = 1.366473
Validation loss decreased (1.391630 --> 1.366473). Saving model ...
-----

Epoch 21
-----
Train Error: Avg loss = 1.319537
Test Error: Accuracy = 51.3%, Avg loss = 1.366601
-----

Epoch 22
-----
Train Error: Avg loss = 1.304318
Test Error: Accuracy = 51.3%, Avg loss = 1.373809
-----

Epoch 23
-----
Train Error: Avg loss = 1.288919
Test Error: Accuracy = 52.2%, Avg loss = 1.347676
Validation loss decreased (1.366473 --> 1.347676). Saving model ...
-----

Epoch 24
-----
Train Error: Avg loss = 1.271301
Test Error: Accuracy = 52.1%, Avg loss = 1.335282
Validation loss decreased (1.347676 --> 1.335282). Saving model ...
-----

Epoch 25
-----
Train Error: Avg loss = 1.258200
Test Error: Accuracy = 52.1%, Avg loss = 1.342228
-----

Epoch 26
-----
Train Error: Avg loss = 1.244255
Test Error: Accuracy = 52.9%, Avg loss = 1.327705
Validation loss decreased (1.335282 --> 1.327705). Saving model ...
-----

Epoch 27
-----
Train Error: Avg loss = 1.229821
```

Test Error: Accuracy = 53.1%, Avg loss = 1.310165
Validation loss decreased (1.327705 --> 1.310165). Saving model ...

Epoch 28

Train Error: Avg loss = 1.212977
Test Error: Accuracy = 52.9%, Avg loss = 1.316234

Epoch 29

Train Error: Avg loss = 1.201382
Test Error: Accuracy = 53.5%, Avg loss = 1.306172
Validation loss decreased (1.310165 --> 1.306172). Saving model ...

Epoch 30

Train Error: Avg loss = 1.188452
Test Error: Accuracy = 54.3%, Avg loss = 1.281335
Validation loss decreased (1.306172 --> 1.281335). Saving model ...

Epoch 31

Train Error: Avg loss = 1.175794
Test Error: Accuracy = 54.7%, Avg loss = 1.282952

Epoch 32

Train Error: Avg loss = 1.159714
Test Error: Accuracy = 55.0%, Avg loss = 1.275567
Validation loss decreased (1.281335 --> 1.275567). Saving model ...

Epoch 33

Train Error: Avg loss = 1.150332
Test Error: Accuracy = 54.2%, Avg loss = 1.292983

Epoch 34

Train Error: Avg loss = 1.137209
Test Error: Accuracy = 54.8%, Avg loss = 1.282207

Epoch 35

Train Error: Avg loss = 1.123812
Test Error: Accuracy = 54.1%, Avg loss = 1.303904

Epoch 36

Train Error: Avg loss = 1.109850
Test Error: Accuracy = 55.0%, Avg loss = 1.270842
Validation loss decreased (1.275567 --> 1.270842). Saving model ...

Epoch 37

Train Error: Avg loss = 1.100081
Test Error: Accuracy = 54.7%, Avg loss = 1.279844

Epoch 38

Train Error: Avg loss = 1.088623
Test Error: Accuracy = 56.0%, Avg loss = 1.261419
Validation loss decreased (1.270842 --> 1.261419). Saving model ...

Epoch 39

Train Error: Avg loss = 1.076307
Test Error: Accuracy = 55.7%, Avg loss = 1.254097
Validation loss decreased (1.261419 --> 1.254097). Saving model ...

Epoch 40

Train Error: Avg loss = 1.063146
Test Error: Accuracy = 56.2%, Avg loss = 1.248550
Validation loss decreased (1.254097 --> 1.248550). Saving model ...

Epoch 41

Train Error: Avg loss = 1.056052
Test Error: Accuracy = 55.3%, Avg loss = 1.276219

Epoch 42

Train Error: Avg loss = 1.044621
Test Error: Accuracy = 55.7%, Avg loss = 1.269160

Epoch 43

Train Error: Avg loss = 1.031281
Test Error: Accuracy = 56.2%, Avg loss = 1.260565

Epoch 44

Train Error: Avg loss = 1.020718
Test Error: Accuracy = 55.7%, Avg loss = 1.270957

Epoch 45

Train Error: Avg loss = 1.009055
Test Error: Accuracy = 56.7%, Avg loss = 1.265475

Epoch 46

```

-----
Train Error: Avg loss = 0.997246
Test Error: Accuracy = 56.4%, Avg loss = 1.255921
-----

Epoch 47
-----
Train Error: Avg loss = 0.990585
Test Error: Accuracy = 56.9%, Avg loss = 1.245865
Validation loss decreased (1.248550 --> 1.245865). Saving model ...
-----

Epoch 48
-----
Train Error: Avg loss = 0.979505
Test Error: Accuracy = 57.4%, Avg loss = 1.245168
Validation loss decreased (1.245865 --> 1.245168). Saving model ...
-----

Epoch 49
-----
Train Error: Avg loss = 0.967405
Test Error: Accuracy = 56.3%, Avg loss = 1.286391
-----

Epoch 50
-----
Train Error: Avg loss = 0.959653
Test Error: Accuracy = 57.2%, Avg loss = 1.235641
Validation loss decreased (1.245168 --> 1.235641). Saving model ...
-----

```

Desempenho no treino e na validação

Gráficos que mostram como foi a redução da loss function de acordo com as épocas de treino, tanto na validação quanto do próprio treino.

```

In [11]: epoch_vec = list(range(1, n_epoch+1))

fig, (ax1, ax2) = plt.subplots(1, 2)
fig.suptitle('Gráficos de Validação')
fig.set_size_inches(12, 5)


ax1.set_title('Perda')
ax1.set_xlabel('Época')
ax1.set_ylabel('Perda')
ax1.plot(epoch_vec, loss_valid_history, label='loss')

ax2.set_title('Acurácia')
ax2.set_xlabel('Época')
ax2.set_ylabel('Acurácia')
ax2.plot(epoch_vec, acc_history, label='accuracy')

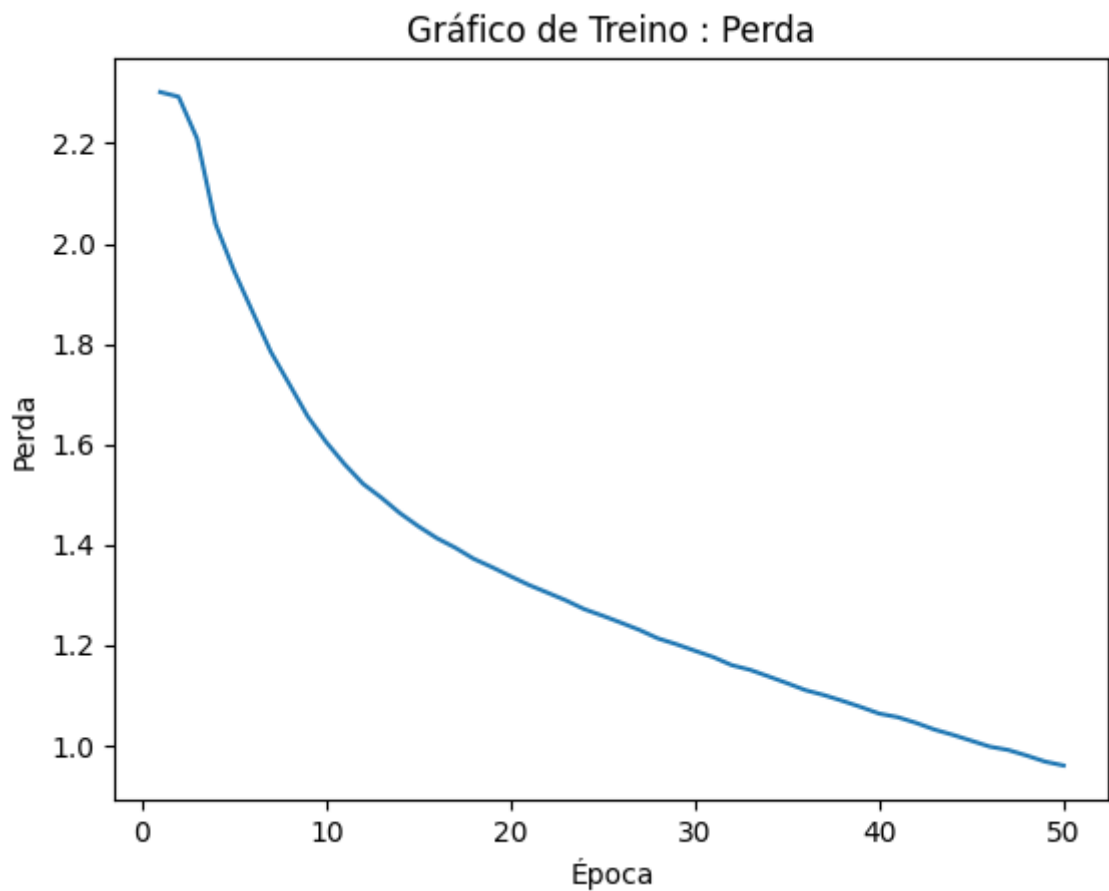
fig.show()

```

C:\Users\sergi\AppData\Local\Temp\ipykernel_15276\2620626654.py:17: UserWarning: Matplotlib is currently using module://matplotlib_inline.backend_inline, which is a non-GUI backend, so cannot show the figure.
fig.show()



```
In [12]: plt.title('Gráfico de Treino : Perda')
plt.plot(epoch_vec, loss_train_history, label='loss')
plt.xlabel('Época')
plt.ylabel('Perda')
plt.show()
```



Testando o Modelo

Pegando o Melhor modelo

```
In [13]: model.load_state_dict(torch.load('models/best_model_trab3.pt'))
```

```
Out[13]: <All keys matched successfully>
```

Função que executa os testes

```
In [14]: def test_loop(dataloader, model):  
    y_pred = []  
    y_true = []  
  
    for X, y in dataloader:  
        pred = model(X)  
        pred = (torch.max(torch.exp(pred), 1)[1]).data.cpu().numpy()  
        y_pred.extend(pred)  
  
        y = y.data.cpu().numpy()  
        y_true.extend(y)  
  
    return y_true, y_pred
```

Resultados

O resultados encontrados são mostrados a seguir, para um problema de classificação com 10 classes uma acurácia acima de 55% deve ser considerada. Além disso, é possível perceber que apenas ao aplicar o modelo da CNN o desempenho entre o trabalho 1 que eu utilizei o mesmo dataset melhorou bastante.

```
In [15]: y_true, y_pred = test_loop(test_loader, model)  
  
cm = confusion_matrix(y_true, y_pred)  
ConfusionMatrixDisplay(cm, display_labels = test_dataset.classes).plot()  
report = classification_report(y_true, y_pred, output_dict=True)  
  
for i in range(len(test_dataset.classes)):  
    print(f"{test_dataset.classes[i]}:\n\tPrecision: {report[str(i)][ 'precision'  
  
print(f"\nAcurácia: {report['accuracy']*100:>.02f}%")  
mcc = matthews_corrcoef(y_true, y_pred)  
print(f"\nMCC: {mcc:>.02f}")
```

airplane:
Precision: 63.40%
Recall: 64.10%
F1-Score: 63.75%

automobile:
Precision: 67.85%
Recall: 70.50%
F1-Score: 69.15%

bird:
Precision: 45.49%
Recall: 51.90%
F1-Score: 48.48%

cat:
Precision: 41.33%
Recall: 39.80%
F1-Score: 40.55%

deer:
Precision: 53.77%
Recall: 46.40%
F1-Score: 49.81%

dog:
Precision: 52.73%
Recall: 37.70%
F1-Score: 43.97%

frog:
Precision: 57.41%
Recall: 70.50%
F1-Score: 63.29%

horse:
Precision: 67.23%
Recall: 63.40%
F1-Score: 65.26%

ship:
Precision: 62.49%
Recall: 76.80%
F1-Score: 68.91%

truck:
Precision: 64.86%
Recall: 56.30%
F1-Score: 60.28%

Acurácia: 57.74%

MCC: 0.53

